

Assignment A03: "Code Rouge" Pitch Document

Version 1.0

Team Members: Josh Zahn, Josh Sourbeer, Brandon Alderman

Instructional Context:

Code Rouge is designed for mid-career medical professionals preparing for their first international humanitarian deployment in Francophone regions of Central or West Africa. Given a simulated high-pressure field hospital environment with limited access to a translator, learners will be able to conduct a basic patient intake in French: establishing rapport, gathering symptoms, and communicating a treatment plan. Learners will thus demonstrate proficiency by accurately triaging the patient within a strict time limit while maintaining a patient Trust Meter score of at least 50%. The game is delivered as a web-based, mobile-friendly application designed for self-directed home use during pre-deployment preparation. This training addresses a documented need within international humanitarian organizations for conversational language competency among field-deployed medical staff (Ärzte ohne Grenzen, n.d.).

High Concept, Theme and Tone:

A deadly outbreak is spreading, and the only thing standing between you and your patients is a language you don't speak. Code Rouge places players in a Francophone field hospital where they learn medical French not from a textbook, but because lives depend on every consultation. The theme is responsibility under pressure. Players feel the discomfort of a competent professional rendered helpless by a language barrier, building through earned fluency toward confidence and connection, and culminating in the high-stakes pressure of a regional outbreak.

Features:

- Speech-to-text AI analysis evaluates pronunciation, accuracy, and cultural appropriateness in real time, providing immediate corrective feedback during patient consultations.
- A persistent Trust Meter rises or decays based on linguistic performance, creating meaningful consequences for every word spoken.
- Trust Meter feeds a cumulative Community Enrichment Score - stronger communication drives more locals to seek treatment, directly reducing outbreak severity.
- A bilingual liaison provides contextual hints at the cost of limited assist tokens, scaffolding player independence over time.
- Five missions of increasing difficulty progress from guided pronunciation to fully autonomous real-time patient intake.

Pedagogical Foundations:

Code Rouge is grounded in the Learning by Doing instructional theory (Schank, Berman, & Macpherson, 1999), which posits that skill development and knowledge acquisition occur best in the context of goals that are relevant and meaningful to learners, and when content is learned through tasks closely related to how it will be used outside the learning environment. The mission places the learner on a realistic NGO deployment in a Francophone region, and the cover story creates an authentic need that motivates learning. Five scenarios of increasing complexity progressively fade scaffolding, from guided pronunciation practice to fully autonomous real-time communication with an AI-powered patient. Resources, including a field guidebook, intake forms, and a bilingual liaison, are well-organized, accessible and provided in context. Feedback is situated and just-in-time through a Trust Meter that responds to the consequences of linguistic accuracy, pronunciation, and cultural appropriateness, supplemented by a post-shift scorecard. This progression of increasingly complex whole tasks is further supported by Merrill's (2013) Task-Centered Principle of Instruction.

Story:

The player takes on the role of a mid-career healthcare professional arriving in the République du Makanda for an NGO medical mission. Upon landing, the player meets an organizational liaison, a bilingual French-English local medical professional, who reveals that most clinic staff and patients speak only French. The liaison eases the player into the language through low-stakes social encounters, starting with low-stakes social encounters to build basic vocabulary. As the player transitions to the clinic, the same liaison serves as their supervising guide, walking them through their first patient intakes and offering translation help and cultural context when needed. Over successive shifts the liaison gradually steps back, and the player takes on consultations independently with increasingly complex patients. The narrative culminates in a regional disease outbreak that demands the player synthesize everything they have learned, interviewing patients under pressure, interpreting medical documents, and piecing together an epidemiological timeline, to contain the crisis before the clinic is overwhelmed.

GamePlay:

During each clinic shift, learners engage in multi-turn consultations where they listen to patients speak French, answer comprehension checks, and respond aloud using speech recognition, or via text input as a fallback. A guidebook and intake forms are available on the desk for reference, and in early missions the liaison provides contextual hints at the cost of limited assist tokens, scaffolding the player's independence over time. Each consultation concludes with a treatment decision whose outcome adjusts a persistent trust meter; strong communication across multiple characters builds a cumulative Community Enrichment score that reflects the player's growing reputation in the region. High community enrichment drives more locals to seek treatment, directly reducing the rate of new outbreaks, meaning the difficulty of the final outbreak scenario is shaped by how well the player communicated throughout the entire game. At the end of every shift, a scorecard breaks performance across five dimensions, tracking cumulative trust, community enrichment, and outbreak rate.

References

Ärzte ohne Grenzen. (n.d.). *Which languages are important for MSF assignments?*
<https://www.aerzte-ohne-grenzen.de/work-with-us/languages>

Hall, R., & Laurence, M. (2023). *Cause and effect game design: A practical toolbox for creating compelling video games*.

Hirumi, A. (2020). *Instructional design principles*. University of Central Florida, Department of Learning Sciences and Educational Research.

Hirumi, A. (2024). *Grounded strategies for facilitating evidence-based design*. University of Central Florida, Department of Learning Sciences and Educational Research.

Merrill, M. D. (2013). *First principles of instruction: Identifying and designing effective, efficient, and engaging instruction*. Pfeiffer.

Schank, R. C., Berman, T. R., & Macpherson, K. A. (1999). Learning by doing. In C. M. Reigeluth (Ed.), *Instructional-design theories and models: A new paradigm of instructional theory* (Vol. 2, pp. 161–181). Lawrence Erlbaum Associates.

GenAI Narrative

Tools Used: Claude Opus via claude.ai (Josh Z), Claude Code via VS Code and mobile (Brandon and Joshua S), Google NotebookLM (Josh Z)

Critical Thinking: Human vs. AI

The most important distinction in this project was recognizing that AI could synthesize and format information across documents, but could not make design judgment calls rooted in course context, team dynamics, or audience awareness. Three examples illustrate this:

First, when selecting a grounded instructional strategy for the Pedagogical Foundations section, AI initially suggested external theories (Vygotsky's Zone of Proximal Development, situated cognition, task-based language teaching). These are defensible choices in isolation, but a human decision was required to recognize that Learning by Doing (Schank, Berman, & Macpherson, 1999) — drawn directly from Professor Hirumi's Grounded Instructional Strategies document — was the strongest fit because it mapped 1:1 to our game's architecture and demonstrated genuine engagement with the assigned course readings. AI cannot distinguish between a "correct" theory and the "best" theory for this specific course.

Second, AI initially produced a Pedagogical Foundations paragraph that spent 30% of its word count defending the theory choice by citing Gagné's taxonomy and Hirumi's strategy selection guidelines. This read like a thesis defense, not a pitch document. The human judgment call was recognizing that the rubric rewards conciseness and application over academic justification — so the entire defensive section was cut, and the paragraph was restructured to show how the theory is applied rather than why it was chosen.

Third, for the Story and Gameplay sections, AI was effective at synthesizing five separate team documents (meeting notes, game goals, gameplay flow, story plan, and the rubric) into coherent draft paragraphs. However, AI could not judge how Brandon's sections should complement — not overlap with — Josh Z's Instructional Context or Josh S's Features. That editorial coordination required human understanding of who wrote what and how the document reads as a whole. AI also struggled with tone calibration: it could not interpret the professor's verbal guidance ("not too verbose") without a human translating that into specific word count targets and stylistic constraints.

A broader complexity worth noting: none of our team members had prior experience with humanitarian medical deployments or French language instruction. This means we were heavily reliant on AI to synthesize domain-specific information (e.g., MSF deployment requirements, field hospital workflows). The risk is that AI-synthesized content can sound authoritative without being accurate. We mitigated this by cross-referencing AI outputs against our interview notes from a practicing pediatric cardiologist who volunteers annually in Tanzania, and by using NotebookLM's grounded retrieval to ensure claims traced back to specific course documents rather than AI-generated assumptions.

Effective (Practical) Use

AI was used iteratively across three workflows, each with distinct prompting strategies:

Workflow 1 — Instructional Context (Josh Z, Claude via claude.ai): The drafting process began by providing Claude with the full assignment rubric, the pitch template, both example pitch documents (Trainien and Mystic Quest), our team's Goal and Learner Analysis, and our teammates' existing drafts. Rather than asking AI to write the section from scratch, the prompt strategy was to draft first, then ask AI to critique against the rubric.

Sample prompt: "Look at this Instructional Context draft. Compare it against the Exceptional criteria in the rubric. Tell me exactly what is missing and how to fix it — do not rewrite it for me."

AI identified that the draft said "pre-deployment preparation" (a time) but never explicitly stated where the learner physically uses the game — a specific rubric requirement. The fix was adding "designed for self-directed home use." AI also flagged that including a raw URL in the paragraph violated the "brief, clear, and concise" standard, prompting a switch to a parenthetical APA citation.

Workflow 2 — Pedagogical Foundations (Josh Z, Claude + NotebookLM): NotebookLM was used to cross-reference the Grounded Instructional Strategies PDF and ID Principles PDF against our game design. The key prompt to Claude was: "Read every word of these two course documents. Find the exact strategy that matches our game's Goal Analysis and tell me why it fits better than alternatives."

Claude mapped all seven events of Learning by Doing to specific Code Rouge features and explained why alternatives (Gagné's Nine Events, Kolb, PBL) were less precise fits. This analysis was then condensed into the final paragraph, with the human decision being what to cut (the Gagné/Hirumi justification) and what to keep (the Merrill closing reference).

Workflow 3 — Story and Gameplay (Brandon, Claude Code): Brandon used Claude Code's ability to read across a project repository, feeding it meeting notes, the game goals document, gameplay flow document, and the story plan simultaneously. Rather than prompting section-by-section, one broad instruction let AI synthesize across all five documents. Brandon then reviewed output against the rubric's Exceptional criteria and adjusted for tone and length.

Evaluation Criteria Used to Refine All AI Outputs:

1. *Rubric alignment* — Does the output address every element the Exceptional descriptor requires? (e.g., Who/What/Where/Why for Instructional Context; theory + citation + application for Pedagogical Foundations)
2. *Conciseness* — Does it stay within the two-page pitch document limit? Is every sentence doing work?
3. *Course grounding* — Are cited theories drawn from course-assigned readings (Hirumi's documents), not just external sources?
4. *Non-redundancy* — Does this section avoid repeating information already covered in a teammate's section?

5. *Tone calibration* — Does it read like a pitch selling a game, not an academic paper defending a theory choice?

Ethical Use

All source material originated from our team's own design discussions, meeting notes, interview data, and course-assigned readings. AI was used to organize, critique, and refine ideas we had already developed — not to generate novel game design concepts. No real patient data, personal information, or proprietary content was used in any prompts. All AI-generated outputs were reviewed, edited, and in several cases rejected before inclusion. One representational consideration: our game depicts healthcare interactions in a fictional Francophone African country. We were careful to ensure AI outputs did not default to stereotypical portrayals of African communities or medical settings, cross-referencing narrative details against our interview notes from an active volunteer in Tanzania. This narrative discloses the full extent of AI involvement across all team members' contributions.